

A decorative graphic on the left side of the slide. It consists of a solid blue parallelogram and a light green parallelogram, both tilted at an angle. The green shape is partially behind the blue one, creating a layered effect. The background of the entire slide is a solid teal color.

# TIME SERIES ANALYSIS STOCK MARKET



# INTRODUCTION


Time series analysis is a powerful statistical tool used to analyze time-ordered data points. In this project, we leverage historical stock price data of Tata Steel to build an ARIMA (AutoRegressive Integrated Moving Average) model. ARIMA is a popular method used in time series forecasting due to its ability to model and predict based on past data points while accounting for different components such as trends and seasonality.

# WHAT IS TIME SERIES ?

## Time Series

→ Data points recorded at specific time intervals to understand underlying patterns.


(1) Components → Trends (Overall direction of the data over time)



→ Seasonality (Repeating pattern of data over a set period of time)



→ Cycle (Repeating but non-Seasonal Patterns in the data.)



→ Variation (Unpredictable ups and downs in the data that cannot be explained by these other components)

# ARIMA MODEL

FORECASTING MODEL

→ ARIMA

Auto Regressive Component

↳ How past values affect future values

I → Integrated (accounts for trends and seasonality)

MA → Moving Average Component  
(Remove non-deterministic or Random movement from time series)

Exponential Smoothing?

forecast time series data that doesn't have a clear trend or seasonality.

# LEARNING IN THIS PROJECT

CHOOSE YOUR STOCK DATA

DATA VISUALIZATION

DATA CLEANING

TIME SERIES PLOT

ACF AND PACF PLOT

STATIONARITY TEST (ADF


DECOMPOSITION

ARIMA MODELING

FORECASTING WITH ARIMA

SARIMA MODELING

FORECAST VS ACTUAL



# START WITH THE CODE

## IMPORT LIBRARIES

```
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller
from sklearn.metrics import mean_squared_error, mean_absolute_error
import statsmodels.api as sm
from statsmodels.tsa.seasonal import seasonal_decompose
import yfinance as yf
import warnings
warnings.filterwarnings("ignore")
```

- **Pandas** is used for data manipulation and analysis.
- **Seaborn** is a visualization library based on matplotlib.
- **Matplotlib** is a plotting library.
- Plot the **Autocorrelation Function (ACF)** and **Partial Autocorrelation Function (PACF)** of a time series.
- **ADFuller** is the Augmented Dickey-Fuller test, used to test the stationarity of a time series.
- Performance metrics used to evaluate the accuracy of the forecast models.
- Statistical models.
- Decompose a time series.
- yFinance is a library that allows for easy downloading of stock market data.
- Use to ignore warnings.



## DOWNLOADING THE DATA AND SAVING IT AS A CSV FILE IN THE GOOGLE COLAB NOTEBOOK

```
[27] # Download historical stock data for Tata Steel from Yahoo Finance
      data = yf.download('TATASTEEL.NS', start='2020-01-01', end='2021-01-01')

      # Save the data to a CSV file
      data.to_csv('tatasteel_stock.csv')
```

## LOADING THE DATA




```
#Read the csv file
df = pd.read_csv('/content/tatasteel_stock.csv')
# Display the first few rows of the DataFrame
print(df.head())
```



	Date	Open	High	Low	Close	Adj Close	\
0	2020-01-01	47.299999	47.650002	46.480000	46.775002	40.974754	
1	2020-01-02	47.200001	48.779999	47.200001	48.485001	42.472713	
2	2020-01-03	48.299999	48.619999	47.945000	48.369999	42.371967	
3	2020-01-06	48.000000	48.000000	47.055000	47.325001	41.456554	
4	2020-01-07	47.549999	48.459999	47.355000	47.610001	41.706207	

	Volume
0	121005300
1	216749610
2	129568630
3	96016080
4	131957880



```
[30] #fill the missing values with forward fill
      if df.isnull().sum().sum() > 0:
          df = df.fillna(method='ffill')
```

## OUTLIER HANDING

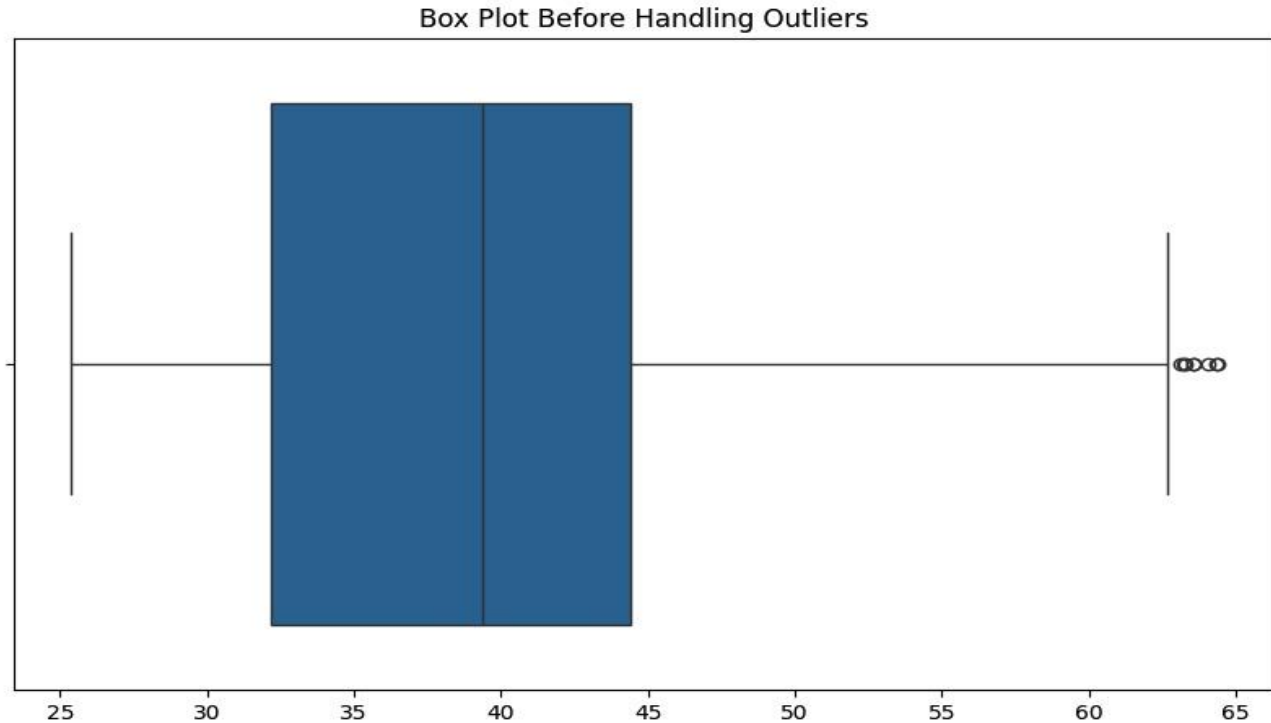
- Detect the outlier.
- Visualize the outlier using BOX PLOT.
- Handle outlier using IQR.



## OUTLIERS IDENTIFICATION



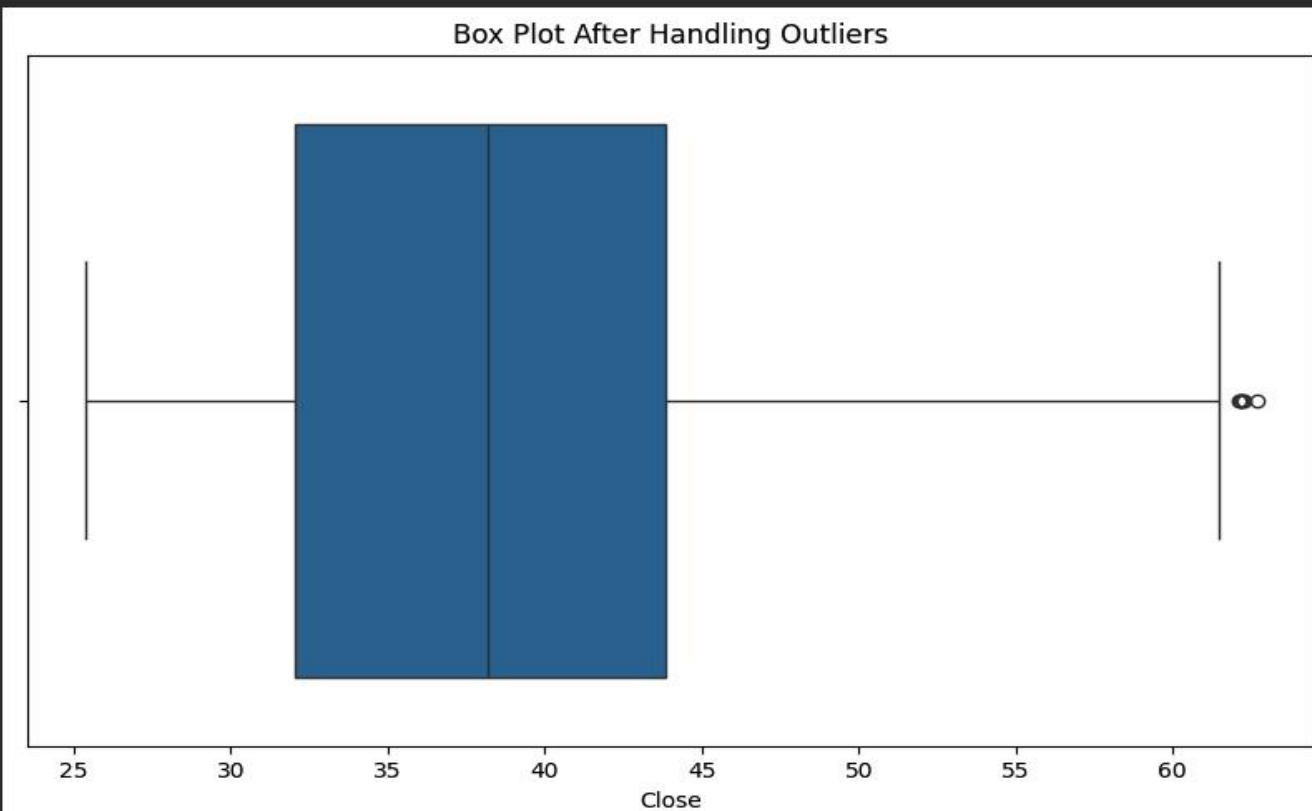
```
#Outlier detection
plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Close'])
plt.title('Box Plot Before Handling Outliers')
plt.show()
```





```
#Calculsting the quartiles
Q1 = df['Close'].quantile(0.25)
Q3 = df['Close'].quantile(0.75)
IQR = Q3 - Q1 #Calculating Interquartile Range
#Calculating upper and lower bound
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
#Filtering out the outliers
df = df[(df['Close'] >= lower_bound) & (df['Close'] <= upper_bound)]
```

```
plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Close'])
plt.title('Box Plot After Handling Outliers')
plt.show()
```




## PLOTTING THE TIME SERIES GRAPH

```
[40] # DATE is containing date and converting it to datetime
      df['Date'] = pd.to_datetime(df['Date'])
      df.set_index('Date', inplace=True)

      # Plot the 'Close' price against the date
      plt.figure(figsize=(10, 6))
      plt.plot(df.index, df['Close'], label='Close Price')
      plt.title('Tata Steel Stock Price')
      plt.xlabel('Date')
      plt.ylabel('Close Price')
      plt.legend()
      plt.show()
```

Tata Steel Stock Price



- Stationarity means that the statistical properties of the time series, such as mean, variance, and autocorrelation, are constant over time.
- Stationary data simplifies the modeling process. Non-stationary data can exhibit trends, seasonality, and other patterns that complicate the model.
- In non-stationary data, the probability distributions can change over time, making statistical inference less reliable.

```
[41] #check the stationarity
def adf_test(series):
    result = adfuller(series)
    print('ADF Statistic: %f' % result[0])
    print('p-value: %f' % result[1])
    for key, value in result[4].items():
        print('\t%s: %.3f' % (key, value))


adf_test(df['Close'])
```

```
⇒ ADF Statistic: -1.313067
   p-value: 0.623173
       1%: -3.459
       5%: -2.874
      10%: -2.573
```

Double-click (or enter) to edit

Since the p-value is greater than 0.05 the time series is non-stationary.





```
[42] # If the data is not stationary, take the first difference
      data_diff = df['Close'].diff().dropna()
      adf_test(data_diff)
```

```
⇒ ADF Statistic: -3.065495
   p-value: 0.029214
       1%: -3.459
       5%: -2.874
      10%: -2.573
```



## Autocorrelation Function (ACF)

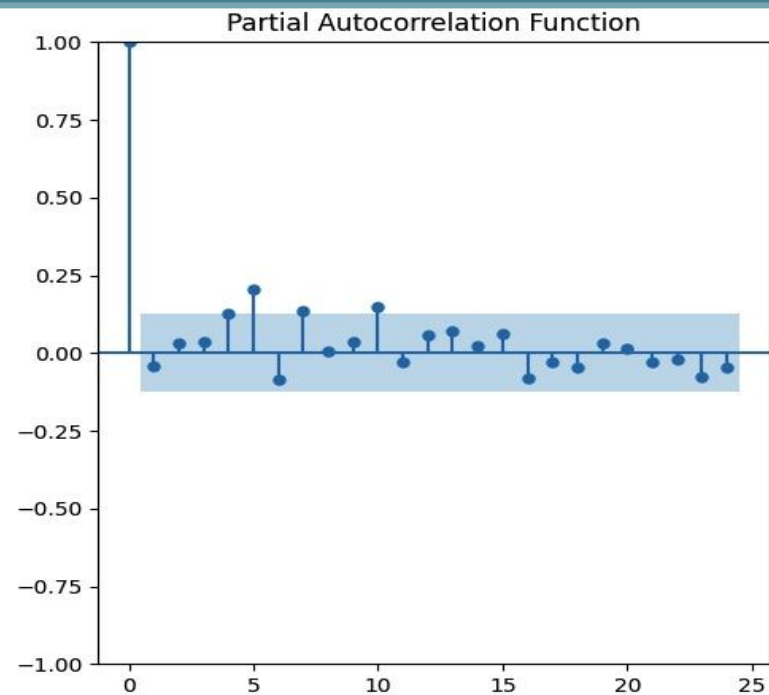

- Quantifies the relationship between the values of the series at different points in time.
- The ACF helps to identify the extent of correlation between the current value and past values of the time series.

## Partial Autocorrelation Function (PACF)

- It essentially isolates the direct relationship between the current value and a lagged value, excluding the influence of other lags.
- The PACF helps to determine the appropriate number of lag terms (autoregressive terms) to include in an ARIMA model.




```
# Plot ACF and PACF
plt.figure(figsize=(12, 6))
plt.subplot(121)
plot_acf(data_diff, ax=plt.gca())
plt.title('Autocorrelation Function')
plt.subplot(122)
plot_pacf(data_diff, ax=plt.gca())
plt.title('Partial Autocorrelation Function')
plt.show()
```



```
[73] # Decompose the time series into trend, seasonal, and residual components
decomposition = seasonal_decompose(df['Close'], model='multiplicative', period=30)
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid


# Plot the decomposition components
plt.figure(figsize=(12, 8))
plt.subplot(411)
plt.plot(df['Close'], label='Original')
plt.legend(loc='upper left')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='upper left')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='upper left')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='upper left')
plt.tight_layout()
plt.show()
```



```
[58] # Split the data into train and test sets
      train_data = data_diff[:-60]
      test_data = data_diff[-60:]

      # Plotting
      plt.figure(figsize=(18, 8))
      plt.grid(True)
      plt.xlabel('Dates', fontsize=20)
      plt.ylabel('Closing Prices', fontsize=20)
      plt.xticks(fontsize=15)
      plt.plot(train_data, 'g', label='Train data', linewidth=2)
      plt.plot(test_data, 'b', label='Test data', linewidth=2)
      plt.legend(fontsize=20, shadow=True, facecolor='lightpink', edgecolor='k')
      plt.show()
```





## GETTING THE BEST ARIMA ORDER

```
import itertools
import statsmodels.api as sm

# Define the p, d, and q parameters to take any value between 0 and 2
p = d = q = range(0, 3)

# Generate all different combinations of p, d, and q triplets
pdq = list(itertools.product(p, d, q))

# Grid search for the optimal ARIMA parameters
best_aic = float("inf")
best_order = None

for order in pdq:
    try:
        model = sm.tsa.ARIMA(train['Close'], order=order)
        results = model.fit()
        aic = results.aic
        if aic < best_aic:
            best_aic = aic
            best_order = order
    except:
        continue

print("Best ARIMA order:", best_order)
print("Best AIC:", best_aic)
```

```
Best ARIMA order: (0, 1, 0)
Best AIC: 587.0052552346217
```



```
history = [x for x in train_data]
# Creating and fitting the ARIMA model
model = ARIMA(history, order=(0, 1, 0))
model_fit = model.fit()
# Displaying the model summary
print(model_fit.summary())
```

```
[61] def train_arima_model(X, y, arima_order):
    # prepare training dataset
    history = [x for x in X]
    predictions = list()

    for t in range(len(y)):
        model = ARIMA(history, order=arima_order)
        model_fit = model.fit()
        yhat = model_fit.forecast()[0] # Get the forecast value
        predictions.append(yhat)
        history.append(y[t])

    # calculate out-of-sample error
    rmse = np.sqrt(mean_squared_error(y, predictions))
    return rmse
```

```
] #model evaluation
def evaluate_models(dataset, test, p_values, d_values, q_values):
    dataset = dataset.astype('float32')
    best_score, best_cfg = float("inf"), None

    for p in p_values:
        for d in d_values:
            for q in q_values:
                order = (p, d, q)
                try:
                    rmse = train_arima_model(dataset, test, order)
                    if rmse < best_score:
                        best_score, best_cfg = rmse, order
                    print('ARIMA%s RMSE=%.3f' % (order, rmse))
                except Exception as e:
                    print(f"Error with order {order}: {e}")
                    continue

    print('Best ARIMA%s RMSE=%.3f' % (best_cfg, best_score))
```



```
# Prepare the training data
history = [x for x in train_data]
predictions = list()
conf_list = list()

for i in range(len(test_data)):
    model = ARIMA(history, order=(0,1,0))
    model_fit = model.fit()
    fc = model_fit.forecast(alpha=0.05) # Get the forecast
    predictions.append(fc[0])
    history.append(test_data[i])

# Convert predictions to a Pandas Series
fc_series = pd.Series(predictions, index=test_data.index)

# Calculate and print RMSE
rmse = np.sqrt(mean_squared_error(test_data, fc_series))
print(f"RMSE is {rmse}")
```



```
RMSE is 1.166065973154891
```

```
[65] # Iterate through the test data
    for t in range(len(test_data)):
        model = SARIMAX(history, order=(0, 1, 0), seasonal_order=(0, 0, 0, 0))
        model_fit = model.fit(dispatch=False)
        fc = model_fit.forecast()
        predictions.append(fc[0])
        history.append(test_data[t])

# Convert predictions to a Pandas Series
fc_series = pd.Series(predictions, index=test_data.index)

# Calculate and print RMSE
rmse = np.sqrt(mean_squared_error(test_data, fc_series))
print('RMSE of SARIMA Model:', rmse)
```




RMSE of SARIMA Model: 1.6484481633429833



## ACTUAL VS PREDICTED GRAPH

```
▶ plt.figure(figsize=(18, 8))  
plt.title('Forecast vs Actual', fontsize=25)  
plt.plot(range(60), predictions, label='Predictions', linewidth=4)  
plt.plot(range(60), test_data, label='Close', linewidth=4)  
plt.legend(fontsize=25, shadow=True, facecolor='lightpink', edgecolor='k')  
plt.show()
```

# Forecast vs Actual





THANK YOU!!